

Umang

LinkedIn: Umang

GitHub: Umang

umangbackend@gmail.com

+91-6395843633

SKILLS

- **Languages** Go, Python, Java, SQL, React
- **Backend & Distributed Systems** Microservices, Event-Driven Architecture, REST APIs, Concurrency, Scalable/High-Throughput Systems, Reliability, Design Review, CI/CD
- **Agent Systems & AI Engineering** LLM Integration, RAG Pipelines, Agentic AI, LangChain, LangGraph, Multi-Step Agent Orchestration, MCP (Tool Servers, Schema Design), Claude Agent SDK, AI Deployment, AI-Assisted Development
- **Data, Cloud & Infrastructure** Kafka, PostgreSQL, TimescaleDB, Redis, Vector Databases, AWS, GCP/Azure, Docker

EXPERIENCE

- **CoinSwitch** Bengaluru, India
Senior Software Engineer **Nov 2025 - Present**
 - **Delegated Access-Based Trading:** Developed a low-latency, **role-based access** system via a unified service, enabling **1M+ active users** to securely grant delegated access to their trading accounts; extended the framework to support authorized access control across **10+** product lines and features.
 - **Reliability, Performance & Incident Management:** Re-architected high-volume trading services for scalability, reducing P95 latency from **380ms to 30ms** for orders flow; modularized core services, expanded automated test coverage, and owned reliability/incident response, cutting MTTR by **60%**.
 - **Real-Time Pricing Pipeline:** Delivered org-wide **INR→USDT** conversion with real-time rate adjustments and hedging-backed position protection, processing **30M+** orders daily.
 - **Observability, AI-Native Development & Debugging:** Shipped **MCP-based AI agents** for on-call/operational workflows; overhauled logging/metrics/alerting for reliability, cutting observability costs from **\$150K to \$10K**.
- **Pixis.ai** Bangalore, India
Senior Software Engineer **Oct 2024 - Oct 2025**
 - **Optimized Core Services Performance:** Reduced memory footprint by **35%** and **P95 latency** by **70%** across critical backend services by eliminating storage inefficiencies, refactoring data models, and migrating graph-heavy workloads to **TimescaleDB**; improved telemetry to streamline debugging.
 - **Refactored High-Concurrency Consumer Service:** Redesigned consumer services using a bounded worker pool model with tuned concurrency limits, eliminating unbounded goroutine growth and improving throughput by **50%** to handle **20K+** messages/sec under burst traffic.
 - **Live Insights & Notifications:** Contributed to live performance-insight and notification features within Prism, surfacing timely campaign alerts that helped marketing teams act on campaign issues **3x** faster.
 - **Cross-Functional Ownership & Mentorship:** Mentored **3** engineers on best practices and code quality, reducing PR review turnaround by **50%**, while owning features from LLD/HLD design through building extensible, reliable systems.
- **LambdaTest** Noida, India
Software Engineer **Jul 2022 - Sep 2024**
 - **Centralized Inventory Microservice Consolidation:** Led the consolidation of distributed inventory services into a single microservice, reducing user launch time by **45%** and eliminating device unavailability through tighter state consistency and service orchestration.
 - **Streamlined Device Onboarding Pipeline:** Established a self-service device integration workflow, replacing manual DB entries and reducing deployment cycle time from **1 week to 1 day (85% improvement)**.
 - **AI-Driven Accessibility Platform:** Designed the end-to-end structure for a scalable, AI-powered accessibility platform and delivered proof-of-concept integrations tailored to enterprise clients' unique use cases, enabling systematic compliance checks across **15+** services and reducing manual audit effort by **60%**.
 - **Growth Initiatives:** Simplified processes and drove strategic improvements resulting in **25%** higher user engagement; created internal tooling (userflow) and ran A/B testing via **feature-flag** implementation.
 - **Micro-Frontend Design & Web Performance:** Built a micro-frontend structure to decouple independently deployable UI modules, reducing page load time by **40%** and improving performance scores; delivered business-critical features across **5+** internal teams, cutting cross-team release cycles by **30%**.

PROJECTS

- **Self-Healing On-Call Agent**
Agentic AI / SRE System **Claude Agent SDK, LangChain, Prometheus, Docker, AI Deployment**
 - **Agentic Triage & Root-Cause Investigation:** Engineered a multi-step reasoning agent using the **Claude Agent SDK** that correlates live Prometheus alerts with logs, traces, and deploy history via custom **MCP** tool servers, cutting manual triage by **70%**.
 - **Tool-Call Design & Guarded Auto-Remediation:** Implemented a propose-approve-execute loop for known failure patterns (restart, scale, rollback) with full action logging and human approval gating, feeding outcomes back into a self-improving runbook store.

ACHIEVEMENTS

- **Google Kick Start:** Ranked **566th** globally among **10,000+** competitive programmers.

EDUCATION

- **Institute of Engineering and Technology** Lucknow, India
B.Tech in Computer Science with a Minor in **Artificial Intelligence** **GPA: 9.0**